



**Department of Computer Science and Engineering (Data Science)**

**Subject: Machine Learning – I (DJ19DSC402)**

**AY: 2022-23**

**Name: Yash Thakar**

**SAP ID.: 60009210205**

Colab Link: [https://colab.research.google.com/drive/1XPme7jyd-ly\\_GtXfdAjB1hYkZQKiwYPQ?usp=sharing](https://colab.research.google.com/drive/1XPme7jyd-ly_GtXfdAjB1hYkZQKiwYPQ?usp=sharing)

**Importing Libraries and Data:**

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

In [2]: path = './data/'
olist_customer = pd.read_csv(path + 'olist_customers_dataset.csv')
olist_geolocation = pd.read_csv(path + 'olist_geolocation_dataset.csv')
olist_order_items = pd.read_csv(path + 'olist_order_items_dataset.csv')
olist_order_payments = pd.read_csv(path + 'olist_order_payments_dataset.csv')
olist_order_reviews = pd.read_csv(path + 'olist_order_reviews_dataset.csv')
olist_orders = pd.read_csv(path + 'olist_orders_dataset.csv')
olist_products = pd.read_csv(path + 'olist_products_dataset.csv')
olist_sellers = pd.read_csv(path + 'olist_sellers_dataset.csv')
olist_translation = pd.read_csv(path + 'product_category_name_translation.csv')
```

**Viewing attributes of all the datasets:**

```
In [3]: dfs = [olist_customer, olist_geolocation, olist_order_items, olist_order_payments, olist_order_reviews, olist_orders, olist_products, olist_sellers, olist_translation]
dfs_names = ['olist_customer', 'olist_geolocation', 'olist_order_items', 'olist_order_payments', 'olist_order_reviews', 'olist_orders', 'olist_products', 'olist_sellers', 'olist_translation']
j = 0
for i in dfs:
    print(dfs_names[j])
    j = j+1
    print(i.shape)
    print(i.columns)
    print()

olist_customer
(99441, 5)
Index(['customer_id', 'customer_unique_id', 'customer_zip_code_prefix', 'customer_city', 'customer_state'], dtype='object')

olist_geolocation
(1000163, 5)
Index(['geolocation_zip_code_prefix', 'geolocation_lat', 'geolocation_lng', 'geolocation_city', 'geolocation_state'], dtype='object')

olist_order_items
(112650, 7)
Index(['order_id', 'order_item_id', 'product_id', 'seller_id', 'shipping_limit_date', 'price', 'freight_value'], dtype='object')

olist_order_payments
(103886, 5)
Index(['order_id', 'payment_sequential', 'payment_type', 'payment_installments', 'payment_value'], dtype='object')

olist_order_reviews
(99224, 7)
Index(['review_id', 'order_id', 'review_score', 'review_comment_title', 'review_comment_message', 'review_creation_date', 'review_answer_timestamp'], dtype='object')

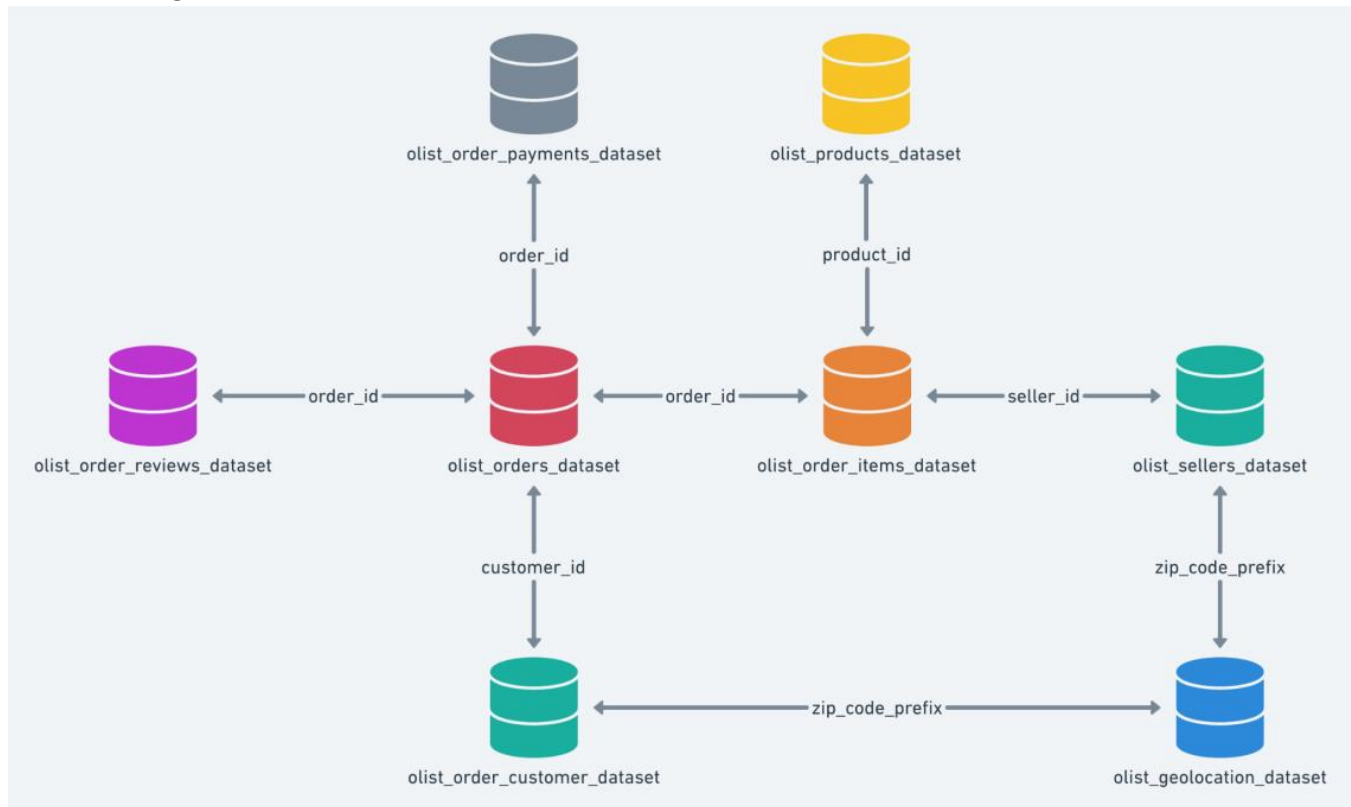
olist_orders
(99441, 8)
Index(['order_id', 'customer_id', 'order_status', 'order_purchase_timestamp', 'order_approved_at', 'order_delivered_carrier_date', 'order_delivered_customer_date', 'order_estimated_delivery_date'], dtype='object')

olist_products
(32951, 9)
Index(['product_id', 'product_category_name', 'product_name_length', 'product_description_length', 'product_photos_qty', 'product_weight_g', 'product_length_cm', 'product_height_cm', 'product_width_cm'], dtype='object')

olist_sellers
(3095, 4)
Index(['seller_id', 'seller_zip_code_prefix', 'seller_city', 'seller_state'], dtype='object')

olist_translation
(71, 2)
Index(['product_category_name', 'product_category_name_english'], dtype='object')
```

## Understanding the DataSchema:



## Number of unique customer ordering on olist:

```
In [4]: olist_customer["customer_unique_id"].nunique()
```

```
Out[4]: 96096
```

## Merging datasets based on the above DataSchema:

```
In [4]: olist_customer["customer_unique_id"].nunique()
```

```
Out[4]: 96096
```

```
In [5]: df = pd.merge(olist_customer, olist_orders, on="customer_id", how='inner')
df = df.merge(olist_order_reviews, on="order_id", how='inner')
df = df.merge(olist_order_items, on="order_id", how='inner')
df = df.merge(olist_products, on="product_id", how='inner')
df = df.merge(olist_order_payments, on="order_id", how='inner')
df = df.merge(olist_sellers, on="seller_id", how='inner')
df = df.merge(olist_translation, on="product_category_name", how='inner')
df.shape
```

```
Out[5]: (115609, 40)
```

```
In [6]: df.head()
```

```
Out[6]:
```

	customer_id	customer_unique_id	customer_zip_code_prefix	customer_city	customer_state
0	06b8999e2fba1a1fbc88172c00ba8bc7	861eff4711a542e4b93843c6dd7feb0	14409	franca	SP
1	8912fc0c3bbf1e2fb35819e21706718	9eae34bbd3a474ec5d07949ca7de67c0	68030	santarem	PA
2	8912fc0c3bbf1e2fb35819e21706718	9eae34bbd3a474ec5d07949ca7de67c0	68030	santarem	PA
3	f0ac8e5a239118859b1734e1087cbb1f	3c799d181c34d51f644bbbc563024db	92480	nova santa rita	RS
4	6bc8d08963a135220ed6c6d098831f84	23397e992b09769fa5e66f9e171a241	25931	mage	RJ

5 rows x 40 columns

Re-assigning appropriate datatypes and Renaming misspelled feature names:

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 115609 entries, 0 to 115608
Data columns (total 40 columns):
#   Column                                  Non-Null Count  Dtype
---  -
0   customer_id                            115609 non-null  object
1   customer_unique_id                     115609 non-null  object
2   customer_zip_code_prefix               115609 non-null  int64
3   customer_city                           115609 non-null  object
4   customer_state                         115609 non-null  object
5   order_id                               115609 non-null  object
6   order_status                           115609 non-null  object
7   order_purchase_timestamp                115609 non-null  object
8   order_approved_at                      115595 non-null  object
9   order_delivered_carrier_date            114414 non-null  object
10  order_delivered_customer_date           113209 non-null  object
11  order_estimated_delivery_date           115609 non-null  object
12  review_id                               115609 non-null  object
13  review_score                            115609 non-null  int64
14  review_comment_title                    13801 non-null  object
15  review_comment_message                  48906 non-null  object
16  review_creation_date                    115609 non-null  object
17  review_answer_timestamp                  115609 non-null  object
18  order_item_id                           115609 non-null  int64
19  product_id                              115609 non-null  object
20  seller_id                               115609 non-null  object
21  shipping_limit_date                     115609 non-null  object
22  price                                   115609 non-null  float64
23  freight_value                           115609 non-null  float64
24  product_category_name                   115609 non-null  object
25  product_name_lenght                     115609 non-null  float64
26  product_description_lenght              115609 non-null  float64
27  product_photos_qty                      115609 non-null  float64
28  product_weight_g                        115608 non-null  float64
29  product_length_cm                       115608 non-null  float64
30  product_height_cm                       115608 non-null  float64
31  product_width_cm                        115608 non-null  float64
32  payment_sequential                       115609 non-null  int64
33  payment_type                             115609 non-null  object
34  payment_installments                    115609 non-null  int64
35  payment_value                           115609 non-null  float64
36  seller_zip_code_prefix                   115609 non-null  int64
37  seller_city                             115609 non-null  object
38  seller_state                             115609 non-null  object
39  product_category_name_english            115609 non-null  object
dtypes: float64(10), int64(6), object(24)
memory usage: 36.2+ MB
```

```
In [8]: for feature in ['order_purchase_timestamp', 'order_approved_at', 'order_delivered_carrier_date',
                        'order_delivered_customer_date', 'order_estimated_delivery_date', 'shipping_limit_date',
                        'review_creation_date', 'review_answer_timestamp']:
        df[feature] = pd.to_datetime(df[feature], errors = 'raise', utc = False)
```

```
In [9]: df.rename(columns = {'product_name_lenght': 'product_name_length',
                             'product_description_lenght': 'product_description_length'}, inplace = True)
```

## Missing Value Analysis and Removal:

We see that 100% of the data is present for most of the important features and the missing data is mostly due to customers choosing not to leave a review for the products they purchased and some failed deliveries. Thus, Dropping the features with high percentage of missing values

```
In [11]: df.shape
```

```
Out[11]: (115609, 40)
```

```
In [12]: Total = df.isnull().sum().sort_values(ascending = False)
Percent = (df.isnull().sum()*100/df.isnull().count()).sort_values(ascending = False)


missing_data = pd.concat([Total, Percent], axis = 1)
print(missing_data)
```

	0	1
review_comment_title	101808	88.062348
review_comment_message	66703	57.697065
order_delivered_customer_date	2400	2.075963
order_delivered_carrier_date	1195	1.033657
order_approved_at	14	0.012110
product_height_cm	1	0.000865
product_weight_g	1	0.000865
product_length_cm	1	0.000865
product_width_cm	1	0.000865
customer_id	0	0.000000
product_name_length	0	0.000000
product_description_length	0	0.000000
product_photos_qty	0	0.000000
payment_sequential	0	0.000000
freight_value	0	0.000000
payment_type	0	0.000000
payment_installments	0	0.000000
payment_value	0	0.000000
seller_zip_code_prefix	0	0.000000
seller_city	0	0.000000
seller_state	0	0.000000
product_category_name	0	0.000000
seller_id	0	0.000000
price	0	0.000000
order_estimated_delivery_date	0	0.000000
customer_zip_code_prefix	0	0.000000
customer_city	0	0.000000
customer_state	0	0.000000
order_id	0	0.000000
order_status	0	0.000000
order_purchase_timestamp	0	0.000000
review_id	0	0.000000
shipping_limit_date	0	0.000000
review_score	0	0.000000
review_creation_date	0	0.000000
review_answer_timestamp	0	0.000000
order_item_id	0	0.000000
product_id	0	0.000000
customer_unique_id	0	0.000000
product_category_name_english	0	0.000000

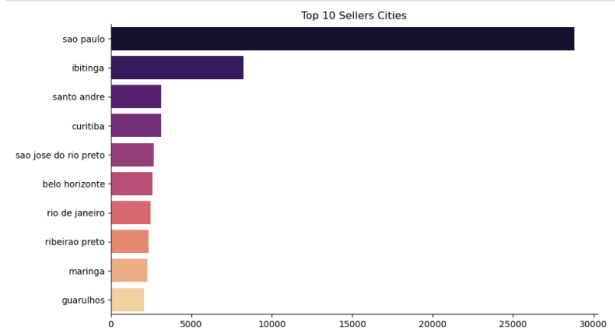
```
In [13]: df.drop(['review_comment_title', 'review_comment_message'], axis = 1, inplace = True)
```

Knowing and understanding your major markets can be a huge deal for e-commerce businesses and thus we plot the top 10 cities that are generating the highest orders

```
In [14]: top_orders_cities = df.groupby("customer_city")["order_id"].count().reset_index().sort_values("order_id", ascending = False)
plt.figure(figsize = (14, 7))
sns.barplot(x = "order_id", y = "customer_city", data = top_orders_cities[:10], palette = 'magma')
plt.xlabel("Number of Orders", fontsize = 14)
plt.ylabel("Cities", fontsize = 14)
plt.yticks(fontsize = 12)
plt.title("Cities Generating the Most Orders (Top 10)", fontsize = 15)
plt.show()
```



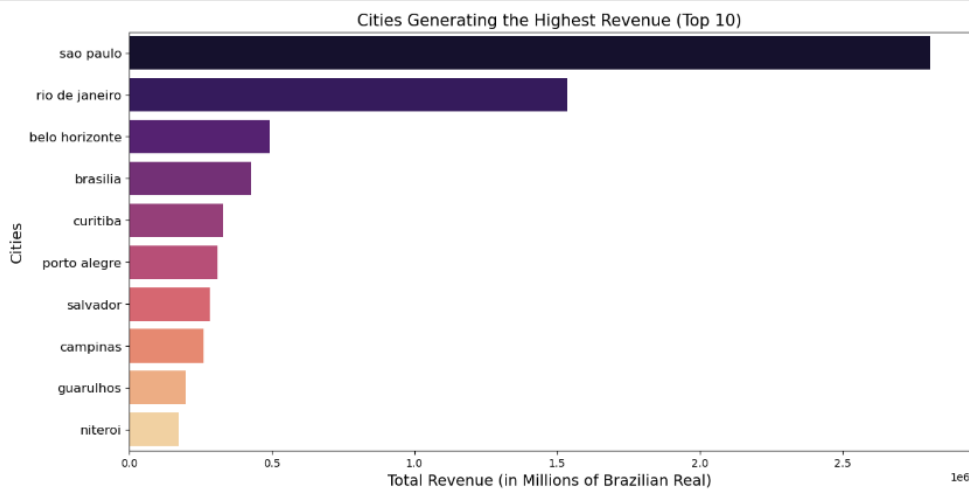
```
In [90]: plt.figure(figsize=(10, 6))
sns.barplot(x = df.seller_city.value_counts().values[:10], y = df.seller_city.value_counts().index[:10], palette= 'magma')
plt.title('Top 10 Sellers Cities')
sns.despine()
```



We can clearly see that most of the orders are coming from Brazil's biggest metropolitan cities - Sao Paulo and Rio de Janeiro, while most sellers are located in sao paulo and ibitinga, this can be a major insight to improve logistical operations.

Similarly lets analyse the cities that generate most revenue

```
In [15]: top_revenue_cities = df.groupby("customer_city")["payment_value"].sum().reset_index().sort_values("payment_value", ascending = False)
plt.figure(figsize = (14, 7))
sns.barplot(x = "payment_value", y = "customer_city", data = top_revenue_cities[:10], palette = 'magma')
plt.xlabel("Total Revenue (in Millions of Brazilian Real)", fontsize = 14)
plt.ylabel("Cities", fontsize = 14)
plt.yticks(fontsize = 12)
plt.title("Cities Generating the Highest Revenue (Top 10)", fontsize = 15)
plt.show()
```



As expected, the cities that generated the most orders, also generated the most revenue.

Understanding Product categories can be of great importance so as to understand which products should the business be focusing on:

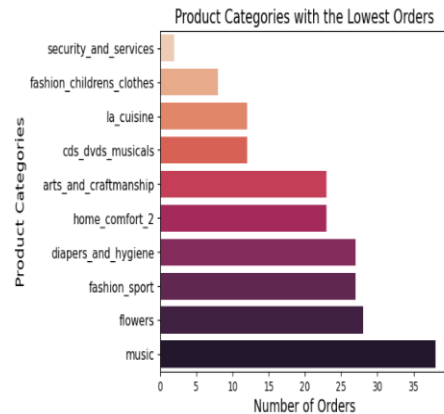
```
In [18]: prodCat_TopOrders = df.groupby(df['product_category_name_english'])['order_id'].nunique().reset_index().sort_values('order_id',
sns.barplot(x = 'order_id', y = 'product_category_name_english', data = prodCat_TopOrders[:10], palette = 'magma')
plt.xlabel("Number of Orders", fontsize = 14)
plt.ylabel("Product Categories", fontsize = 14)
plt.tick_params(axis = 'y', labelsize = 12)
plt.title("Product Categories with the Highest Orders (Top 10)", fontsize = 15)
```

Out[18]: Text(0.5, 1.0, 'Product Categories with the Highest Orders (Top 10)')



```
In [52]: prodCat_TopOrders = df.groupby(df['product_category_name_english'])['order_id'].nunique().reset_index().sort_values('order_id',
sns.barplot(x = 'order_id', y = 'product_category_name_english', data = prodCat_TopOrders[:10], palette = 'rocket_r')
plt.xlabel("Number of Orders", fontsize = 14)
plt.ylabel("Product Categories", fontsize = 14)
plt.tick_params(axis = 'y', labelsize = 12)
plt.title("Product Categories with the Lowest Orders", fontsize = 15)
```

Out[52]: Text(0.5, 1.0, 'Product Categories with the Lowest Orders')



The best Product category are bed\_bath\_table and the worst is security\_and\_services.

But since here many categories can be grouped into one major category to simplify our understanding we do so manually.

```
In [20]: def classify_cat(x):

    if x in ['office_furniture', 'furniture_decor', 'furniture_living_room', 'kitchen_dining_laundry_garden_furniture', 'bed_bath_table']:
        return 'Furniture'

    elif x in ['auto', 'computers_accessories', 'musical_instruments', 'consoles_games', 'watches_gifts', 'air_conditioning', 'telephony']:
        return 'Electronics'

    elif x in ['fashion_female_clothing', 'fashion_male_clothing', 'fashion_bags_accessories', 'fashion_shoes', 'fashion_sport']:
        return 'Fashion'

    elif x in ['housewares', 'home_comfort', 'home_appliances', 'home_appliances_2', 'flowers', 'construction_tools_garden', 'garden_tools']:
        return 'Home & Garden'

    elif x in ['sports_leisure', 'toys', 'cds_dvds_musicals', 'music', 'dvds_blu_ray', 'cine_photo', 'party_supplies', 'christmas']:
        return 'Entertainment'

    elif x in ['health_beauty', 'perfumery', 'diapers_and_hygiene']:
        return 'Beauty & Health'

    elif x in ['food_drink', 'drinks', 'food']:
        return 'Food & Drinks'

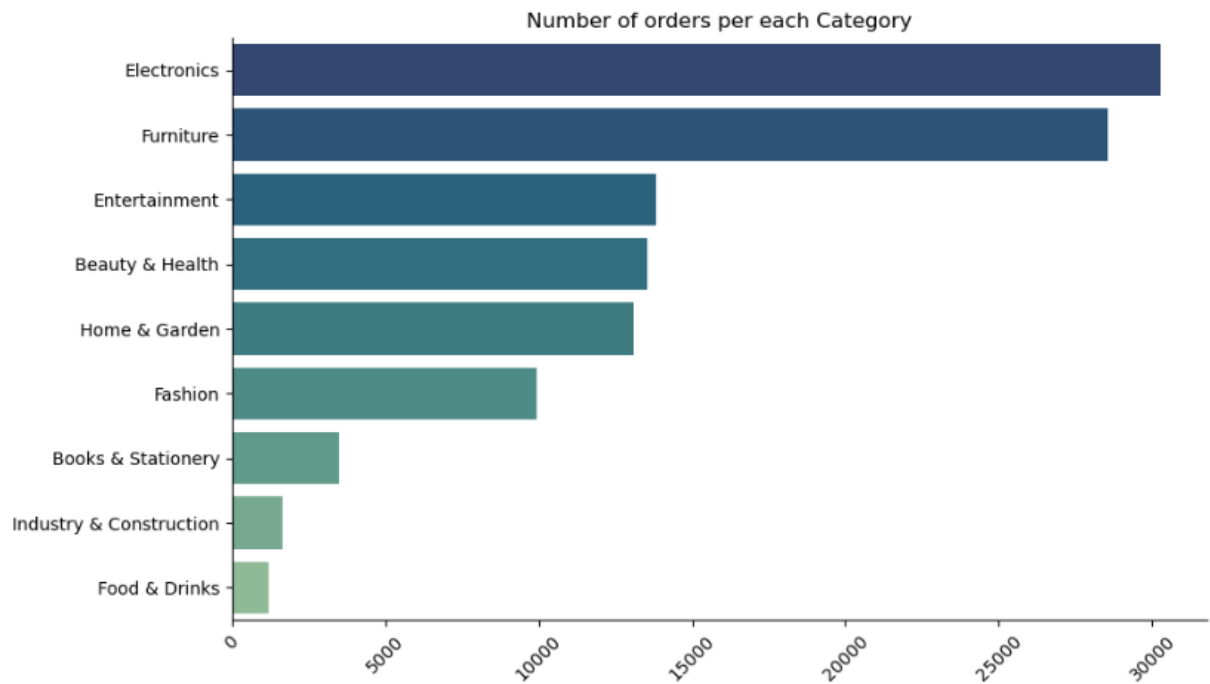
    elif x in ['books_general_interest', 'books_technical', 'books_imported', 'stationery']:
        return 'Books & Stationery'

    elif x in ['construction_tools_construction', 'construction_tools_safety', 'industry_commerce_and_business', 'agro_industry']:
        return 'Industry & Construction'

df['product_category'] = df.product_category_name_english.apply(classify_cat)
```

Now we can observe in the below bar chart that the Electronics is the best performing category followed by Furniture while Foods and Drinks have the least amount of orders on the E-commerce platform

```
In [21]: plt.figure(figsize=[10, 6])
sns.barplot(x = df.product_category.value_counts().values, y = df.product_category.value_counts().index, palette= 'crest_r')
plt.title('Number of orders per each Category')
plt.xticks(rotation = 45)
sns.despine()
```




Analysing times at which the customer is likely to order from the platform

```
In [16]: clrp = sns.color_palette("hls", 1)

orders_byHour = df.groupby(df.order_purchase_timestamp.dt.hour)['order_id'].nunique().reset_index()
plt.figure(figsize = (15, 5))
sns.barplot(x = 'order_purchase_timestamp', y = 'order_id', data = orders_byHour, palette = clrp)
plt.xlabel("Hour of Day", fontsize = 14)
plt.ylabel("Number of Orders", fontsize = 14)
plt.title("Orders by Hour", fontsize = 15)
plt.show()
```



```
In [17]: orders_byDays = df.groupby(df.order_purchase_timestamp.dt.day_name())['order_id'].nunique().reset_index().sort_values('order_id')
plt.figure(figsize = (10, 4))
sns.barplot(x = 'order_purchase_timestamp', y = 'order_id', data = orders_byDays, palette = 'magma')
plt.xlabel("Day of Week", fontsize = 14)
plt.ylabel("Number of Orders", fontsize = 14)
plt.title("Orders by Day of Week", fontsize = 15)
plt.show()
```



We observe that the amount of orders pouring in on Monday to Wednesday is significantly larger than the remaining half of the week.

This can be a useful insight when arranging the logistics for delivery.

And the hourly engagement of orders can be useful if the E-commerce platform focuses on increasing its Food&Drink Product Category by providing one-day or faster deliveries



While we are on the matter of delivery lets take into account reviews and Delivery times, as these are some quantitative factors that affect customer satisfaction:

```
In [24]: # Distribution of review scores
plt.figure(figsize=(8, 6))
sns.countplot(x='review_score', data=df)
plt.title('Distribution of Review Scores', fontsize=16)
plt.xlabel('Review Score', fontsize=14)
plt.ylabel('Count', fontsize=14)
plt.tick_params(axis='both', which='major', labelsize=12)
sns.despine()
plt.show()
```



We can observe that customers are mostly satisfied as there are large number of 4 and 5 Star reviews

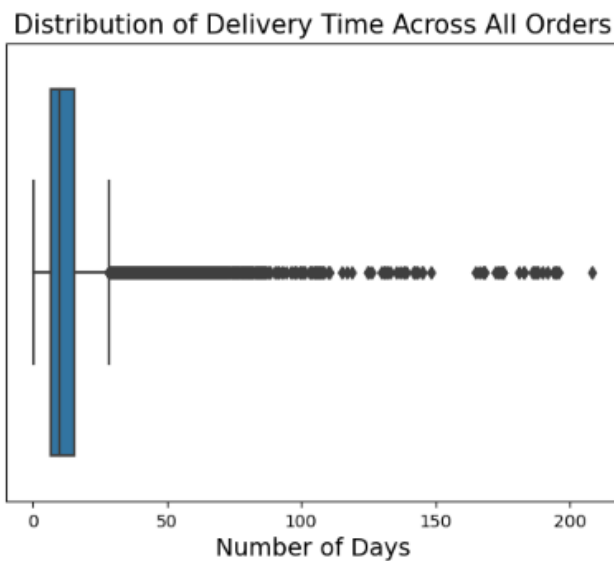
To measure Delivery times,

```
In [28]: deliveryTime = (df["order_delivered_customer_date"] - df["order_purchase_timestamp"])
deliveryTime_Seconds = deliveryTime.apply(lambda x: x.total_seconds())
df['deliveryTime_Days'] = round(deliveryTime_Seconds/86400, 2)
```

```
In [29]: df['deliveryTime_Days'].describe()
```

```
Out[29]: count    113209.000000
mean         12.442129
std           9.356006
min           0.530000
25%           6.740000
50%          10.190000
75%          15.500000
max          208.350000
Name: deliveryTime_Days, dtype: float64
```

```
In [30]: sns.boxplot(df.deliveryTime_Days, orient = 'h', showfliers = True)
plt.xlabel("Number of Days", fontsize = 14)
plt.yticks([])
plt.title('Distribution of Delivery Time Across All Orders', fontsize = 15)
plt.show()
```



We can come to the conclusion from the Box and Whisker Plot, that on an average a delivery takes 12 days to reach the customer, which is significantly slower by modern standards.

We can also observe a huge amount of Outliers in the Box Plot pointing towards inconsistent Delivery Times

Plotting Average arrival / Delivery times for each product Category:

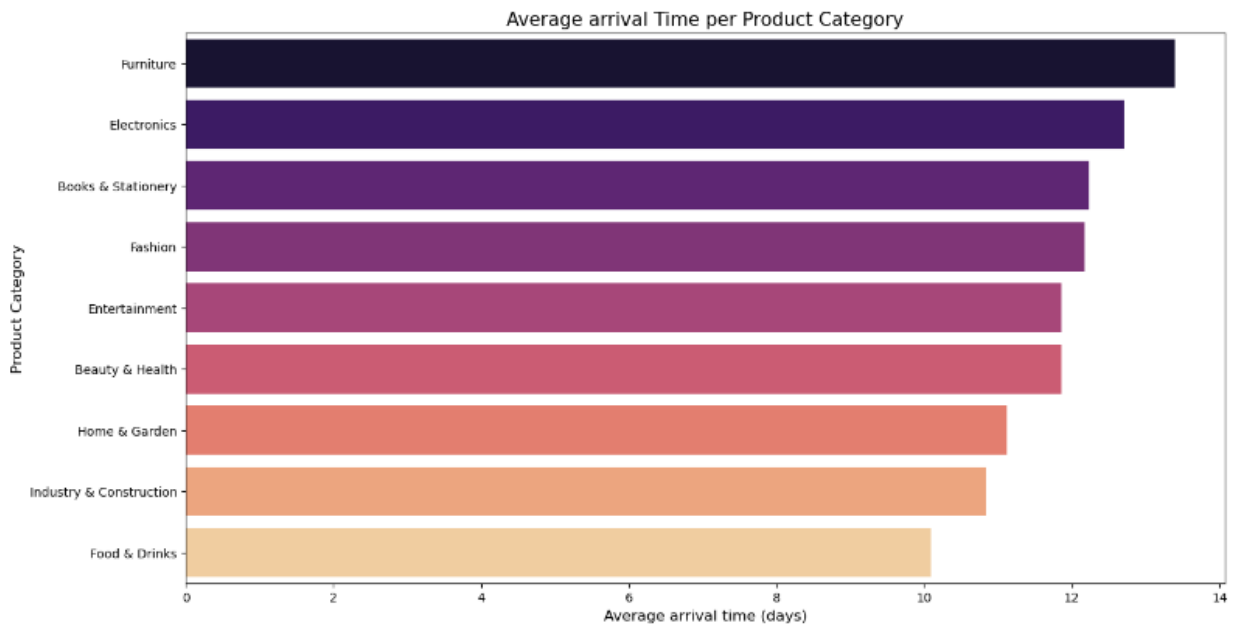
```
In [86]: df['arrival_days'] = (df['order_delivered_customer_date'].dt.date - df['order_purchase_timestamp'].dt.date).dt.days

# Group product category by average arrival time
ship_per_cat = df.groupby('product_category')[['arrival_days']].mean().sort_values(by='arrival_days', ascending=False)
ship_per_cat.reset_index(inplace=True)

plt.figure(figsize=[15, 8])
sns.barplot(x = ship_per_cat.arrival_days, y= ship_per_cat.product_category, palette= 'magma')

plt.title('Average arrival Time per Product Category', fontsize= 15)
plt.xlabel('Average arrival time (days)', fontsize= 12)
plt.ylabel('Product Category', fontsize= 12)
```

Out[86]: Text(0, 0.5, 'Product Category')




We can observe that the Best Performing Categories of Furniture and Electronics have Delivery times equal to or greater than the average, this could be a huge drawback going further as many of the customers will be experiencing slow delivery times.

Also to increase orders and thus revenue from the Food&Drinks Category we need to significantly shorten delivery durations to deliver foods that can be perceived fresh by the customer,

Also faster delivery times in this category would mean customers will more likely to be order Food & Drinks to satisfy their impulsive cravings, thus boosting revenue.


Now lets take a look at how Delivery time affect Review Scores:

```
In [33]: sns.boxplot(x = "review_score", y = "deliveryTime_Days", data = df, showfliers = False, palette = 'rocket')
plt.xlabel("Review Score", fontsize = 14)
plt.ylabel("Delivery Time (in Days)", fontsize = 14)
plt.title("Delivery Time vs. Review Scores", fontsize = 15)
plt.show()
```



As observed in the above Box Plot, there's a slight correlation between delivery times and review scores. The longer it takes for an order to be delivered, the more likely it is to receive a low review score.

```
In [78]: # Contingency table of order status vs. review score
cont_table = pd.crosstab(df['order_status'], df['review_score'])
plt.figure(figsize=(8, 6))
sns.heatmap(cont_table, cmap='Blues', annot=True, fmt='d')
plt.title('Order Status vs. Review Score', fontsize=16)
plt.xlabel('Review Score', fontsize=14)
plt.ylabel('Order Status', fontsize=14)
plt.tick_params(axis='both', which='major', labelsize=12)
sns.despine()
plt.show()
```




We can observe that the proportion of negative reviews for a cancelled order or a order that shows shipped/invoiced/processing as its order status is significantly high and is a sign of unsatisfied customers.

On classifying the delivers into On-time and Late we can observe that 91.7% of deliveries only make it on time, which is a number that needs improvement to improve customer retention:

```
In [34]: df['seller_to_carrier_status'] = (df['shipping_limit_date'].dt.date - df['order_delivered_carrier_date'].dt.date).dt.days
# Now calssify the duration into 'OnTime/Early' & 'Late'
df['seller_to_carrier_status'] = df['seller_to_carrier_status'].apply(lambda x : 'OnTime/Early' if x >=0 else 'Late')


In [35]: # First get difference between estimated delivery date and actual delivery date in days
df['arrival_status'] = (df['order_estimated_delivery_date'].dt.date - df['order_delivered_customer_date'].dt.date).dt.days
# Now Classify the duration in 'OnTime/Early' & 'Late'
df['arrival_status'] = df['arrival_status'].apply(lambda x : 'OnTime/Early' if x >=0 else 'Late')

In [36]: plt.figure(figsize=[30,8])
Values = df.arrival_status.value_counts().values
Labels = df.arrival_status.value_counts().index
plt.pie(Values, explode=(0.05, 0.1), labels= ['OnTime/Early', 'Late'], autopct='%1.1f%%')
plt.show()
```



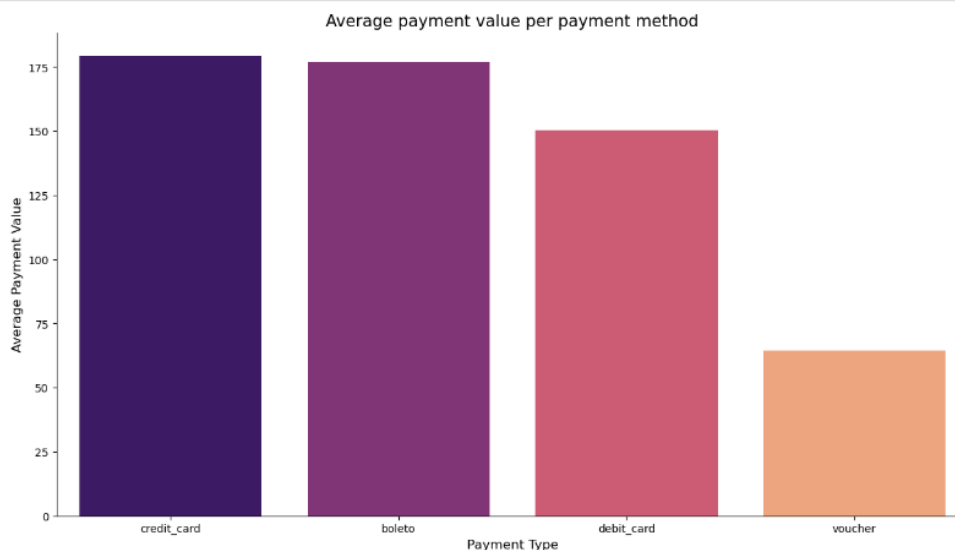
The e-commerce platform has multiple payment methods available, the way in which customers pay can indicate customer preferences and point out customer traits:

```
In [75]: plt.figure(figsize = (10, 6.5))
payment_type_counts = df.payment_type.value_counts()
plt.pie(x = payment_type_counts.to_list(), labels = ['Credit Card', 'Boleto', 'Voucher', 'Debit Card'],
        autopct = '%1.2f%%', explode = (0, 0.025, 0.05, 0.2), colors = sns.color_palette('rocket_r'),
        textprops = {'fontsize': 12})
plt.title("Payment Methods", fontsize = 20)
plt.show()
```



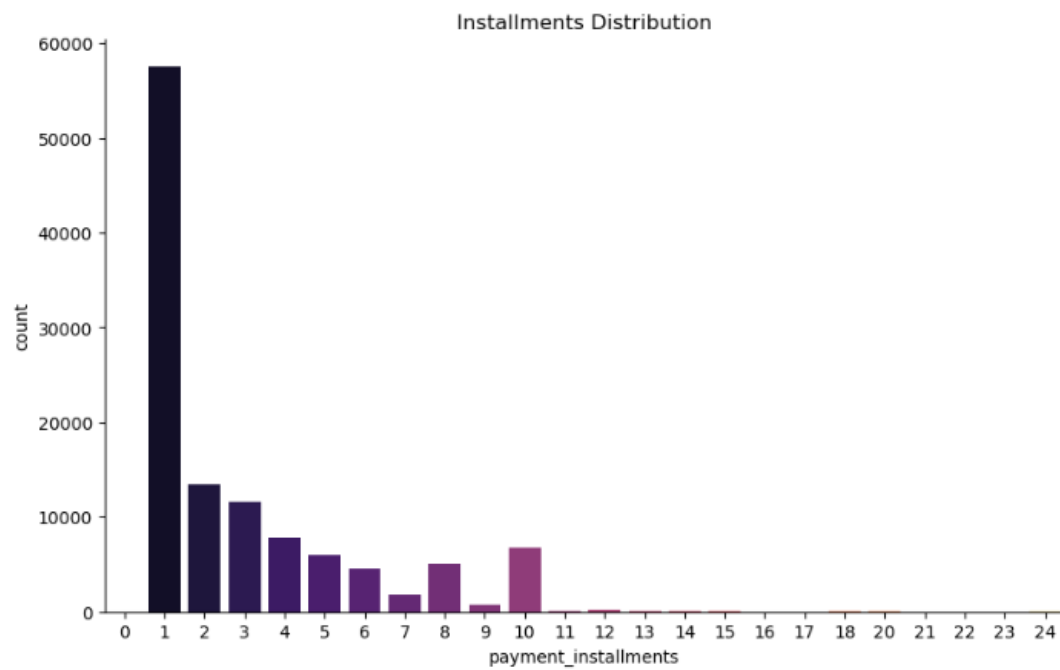
```
In [87]: # Group each payment type by average payment value
payment_methods = df.groupby('payment_type')[['payment_value']].mean().sort_values(by='payment_value', ascending=False)
payment_methods.reset_index(inplace=True)

# Plot Average payments per payment method
plt.figure(figsize=[15, 8])
sns.barplot(x = payment_methods.payment_type, y = payment_methods.payment_value, palette='magma')
plt.title('Average payment value per payment method', fontsize= 15)
plt.xlabel('Payment Type', fontsize= 12)
plt.ylabel('Average Payment Value', fontsize= 12)
sns.despine()
```



We can observe that the even though a significantly more amount of customers pay via credit\_card over Boleto, the average payment value per payment method is approximately the same for both but credit\_card fees that are payable to the service providers are likely to eat away a chunk of the profits.

```
In [79]: plt.figure(figsize=[10, 6])
sns.countplot(x = df.payment_installments, palette= 'magma')
plt.title('Installments Distribution')
sns.despine()
```



We can observe here that a significant number of customer pay in installments ranging from 1 – 12




## Analysing Revenue and Sales:

```
In [40]: monthly_sales = df.resample('M', on='order_purchase_timestamp')['order_item_id'].count()
monthly_sales = monthly_sales['2016-01':'2018-05']

plt.figure(figsize=(10,6))
plt.plot(monthly_sales.index, monthly_sales.values, '-o', color='purple')

plt.title('Sales', fontsize=16)
plt.xlabel('Date', fontsize=14)
plt.ylabel('Number of Products Sold', fontsize=14)
plt.grid(True)
```



There was a huge growth in the sales from 2016 to 2017, hitting its peak in the end of 2017 after which the sales fell down and became inconsistent.


With a major hit in sales observed in the start of a 2018 right after the company made record sales.

A similar trend can be observed in the Revenue:

```
In [50]: monthly_revenue = df['payment_value'].resample('M').sum()
monthly_revenue = monthly_revenue['2016-01': '2018-05']

plt.figure(figsize=(10,6))
plt.plot(monthly_revenue.index, monthly_revenue.values, '-o', color='purple')

plt.title('Revenue', fontsize=16)
plt.xlabel('Year-Month', fontsize=14)
plt.ylabel('Revenue (BRL)', fontsize=14)
plt.grid(True)
plt.show()
```



Although Revenue is constantly growing after the dip in the beginning of 2018/04 to 2018/06 in spite of a sales being a little low throughout this period, this might indicate organizational efficiency and improved margins in the business.