



Department of Computer Science and Engineering (Data Science)

Subject: Machine Learning – I (DJ19DSC402)

AY: 2022-23

Name: Yash Thakar

SAP ID.: 60009210205

Colab Links:

Customer Segmentation and Customer Churn Classification:

<https://colab.research.google.com/drive/1e4k24Q2WofuHN8ykUimO9jBuBHRqbgd8>

Sales Prediction:

<https://colab.research.google.com/drive/1FtT0cTZpG4SUipMQbaPrn1kjfS8wiq02>

Customer Segmentation and Customer Churn Classification

Importing Necessary Libraries and processed dataset from previously done EDA

```
import numpy as np
from datetime import date, datetime, timedelta
import pandas as pd
import matplotlib.pyplot as plt

import plotly.graph_objects as go
import plotly.express as px
from plotly.subplots import make_subplots

from plotly.offline import init_notebook_mode, iplot
init_notebook_mode(connected=True)

df = pd.read_csv(path+'merged_data.csv')

In [46]: df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 115609 entries, 0 to 115608
Data columns (total 39 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   Unnamed: 0          115609 non-null  int64
 1   customer_id         115609 non-null  object
 2   customer_unique_id  115609 non-null  object
 3   customer_zip_code_prefix  115609 non-null  int64
 4   customer_city       115609 non-null  object
 5   customer_state      115609 non-null  object
 6   order_id            115609 non-null  object
 7   order_status        115609 non-null  object
 8   order_purchase_timestamp  115609 non-null  datetime64[ns]
 9   order_approved_at   115595 non-null  datetime64[ns]
10   order_delivered_carrier_date  114414 non-null  datetime64[ns]
11   order_delivered_customer_date  113209 non-null  datetime64[ns]
12   order_estimated_delivery_date  115609 non-null  datetime64[ns]
13   review_id           115609 non-null  object
14   review_score        115609 non-null  int64
15   review_creation_date  115609 non-null  datetime64[ns]
16   review_answer_timestamp  115609 non-null  datetime64[ns]
17   order_item_id       115609 non-null  int64
18   product_id          115609 non-null  object
19   seller_id           115609 non-null  object
20   shipping_limit_date  115609 non-null  datetime64[ns]
21   price               115609 non-null  float64
22   freight_value       115609 non-null  float64
23   product_category_name  115609 non-null  object
24   product_name_length  115609 non-null  float64
25   product_description_length  115609 non-null  float64
26   product_photos_qty   115609 non-null  float64
27   product_weight_g     115608 non-null  float64
28   product_length_cm    115608 non-null  float64
29   product_height_cm    115608 non-null  float64
30   product_width_cm     115608 non-null  float64
31   payment_sequential   115609 non-null  int64
32   payment_type         115609 non-null  object
33   payment_installments  115609 non-null  int64
34   payment_value        115609 non-null  float64
35   seller_zip_code_prefix  115609 non-null  int64
36   seller_city          115609 non-null  object
37   seller_state         115609 non-null  object
38   product_category_name_english  115609 non-null  object
dtypes: datetime64[ns](8), float64(10), int64(7), object(14)
memory usage: 34.4+ MB
```

RFM analysis:

RFM analysis is a marketing technique used to quantitatively rank and group customers based on the recency, frequency and monetary total of their recent transactions to identify the best customers and perform targeted marketing campaigns.

RFM analysis ranks each customer on the following factors:

- **Recency:** How recent was the customer's last purchase? Customers who recently made a purchase will still have the product on their mind and are more likely to purchase or use the product again. Businesses often measure recency in days. But, depending on the product, they may measure it in years, weeks or even hours.
- **Frequency:** How often did this customer make a purchase in a given period? Customers who purchased once are often are more likely to purchase again. Additionally, first time customers may be good targets for follow-up advertising to convert them into more frequent customers.
- **Monetary:** How much money did the customer spend in a given period? Customers who spend a lot of money are more likely to spend money in the future and have a high value to a business.

Recency

```
In [5]: recency = df.groupby('customer_unique_id', as_index=False)['order_purchase_timestamp'].max()
recency.rename(columns={'order_purchase_timestamp': 'LastPurchaseDate'}, inplace = True)
recency.head()
```

```
Out[5]:
```

	customer_unique_id	LastPurchaseDate
0	0000366f3b9a7992bf8c76cdf3221e2	2018-05-10 10:56:27
1	0000b849f77a49e4a4ce2b2a4ca5be3f	2018-05-07 11:11:27
2	0000f46a3911fa3c0805444483337064	2017-03-10 21:05:03
3	0000f8ccb0745a8a4b88665a16c9f078	2017-10-12 20:29:41
4	0004aac84e0df4da2b147fca70cf8255	2017-11-14 19:45:42

```
In [6]: recent_date = df['order_purchase_timestamp'].dt.date.max()
print('The last recent date in the available dataset is: ', recent_date)
```

The last recent date in the available dataset is: 2018-09-03

```
In [7]: recency['Recency'] = recency['LastPurchaseDate'].dt.date.apply(lambda x: (recent_date - x).days)
recency.head()
```

```
Out[7]:
```

	customer_unique_id	LastPurchaseDate	Recency
0	0000366f3b9a7992bf8c76cdf3221e2	2018-05-10 10:56:27	116
1	0000b849f77a49e4a4ce2b2a4ca5be3f	2018-05-07 11:11:27	119
2	0000f46a3911fa3c0805444483337064	2017-03-10 21:05:03	542
3	0000f8ccb0745a8a4b88665a16c9f078	2017-10-12 20:29:41	326
4	0004aac84e0df4da2b147fca70cf8255	2017-11-14 19:45:42	293

Frequency

```
In [8]: frequency = df.groupby(["customer_unique_id"]).agg({"order_id": "nunique").reset_index()
frequency.rename(columns={"order_id": "Frequency"}, inplace=True)
frequency.head()
```

Out[8]:

	customer_unique_id	Frequency
0	0000366f3b9a7992bf8c76cfd3221e2	1
1	0000b849f77a49e4a4ce2b2a4ca5be3f	1
2	0000f46a3911fa3c0805444483337064	1
3	0000f6ccb0745a6a4b88665a16c9f078	1
4	0004aac84e0df4da2b147fca70cf8255	1

Monetary

```
In [9]: monetary = df.groupby('customer_unique_id', as_index=False)['payment_value'].sum()
monetary.rename(columns={"payment_value": "Monetary"}, inplace=True)
monetary.head()
```

Out[9]:

	customer_unique_id	Monetary
0	0000366f3b9a7992bf8c76cfd3221e2	141.90
1	0000b849f77a49e4a4ce2b2a4ca5be3f	27.19
2	0000f46a3911fa3c0805444483337064	86.22
3	0000f6ccb0745a6a4b88665a16c9f078	43.62
4	0004aac84e0df4da2b147fca70cf8255	196.89

Merging RFM

```
In [10]: rfm = recency.merge(frequency, on='customer_unique_id')
rfm = rfm.merge(monetary, on='customer_unique_id').drop(columns='LastPurchaseDate')
rfm.head()
```

Out[10]:

	customer_unique_id	Recency	Frequency	Monetary
0	0000366f3b9a7992bf8c76cfd3221e2	116	1	141.90
1	0000b849f77a49e4a4ce2b2a4ca5be3f	119	1	27.19
2	0000f46a3911fa3c0805444483337064	542	1	86.22
3	0000f6ccb0745a6a4b88665a16c9f078	326	1	43.62
4	0004aac84e0df4da2b147fca70cf8255	293	1	196.89

Labelling RFM (to get scores):

Recency Labels

```
ll_r = rfm.Recency.quantile(0.25)
mid_r = rfm.Recency.quantile(0.50)
ul_r = rfm.Recency.quantile(0.75)
print(ll_r, mid_r, ul_r)
```

```
119.0 223.0 352.0
```

```
def recency_label(recent):
    if recent <= ll_r:
        return 1
    elif (recent > ll_r) and (recent <= mid_r):
        return 2
    elif (recent > mid_r) and (recent <= ul_r):
        return 3
    elif recent > ul_r:
        return 4
```

```
rfm['recency_label'] = rfm.Recency.apply(recency_label)
rfm.head()
```

	customer_unique_id	Recency	Frequency	Monetary	recency_label
0	0000366f3b9a7992bf8c76cfd3221e2	116	1	141.90	1
1	0000b849f77a49e4a4ce2b2a4ca5be3f	119	1	27.19	1
2	0000f46a3911fa3c0805444483337064	542	1	86.22	4
3	0000f6ccb0745a6a4b88665a16c9f078	326	1	43.62	3
4	0004aac84e0df4da2b147fca70cf8255	293	1	196.89	3

Recency label breakdown:

1 - These are the customers who whose visit date(s) are the most recent. (Recency value within the 25% quantile)

2 - These are the customers who whose visit date(s) are not very recent. (Recency value between 25% and 50% quantile)

3 - These are the customers who whose visit date(s) are somewhat recent. (Recency value between 50% and 75% quantile)

4 - These are the customers who whose visit date(s) are the oldest. (Recency value more than 75% quantile)

Frequency Lables:

```
rfm.Frequency.value_counts()
```

```
1    90589
2    2581
3     179
4      30
5       9
6       3
7       3
9       1
15      1
Name: Frequency, dtype: int64
```

```
def frequency_label(frequent):
    if frequent == 1:
        return 4
    elif frequent == 2:
        return 3
    elif frequent == 3:
        return 2
    elif frequent > 3:
        return 1
```

```
rfm['frequency_label'] = rfm.Frequency.apply(frequency_label)
rfm.head()
```

	customer_unique_id	Recency	Frequency	Monetary	recency_label	frequency_label
0	0000366f3b9a7992bf8c76cfd3221e2	116	1	141.90	1	4
1	0000b849f77a49e4a4ce2b2a4ca5be3f	119	1	27.19	1	4
2	0000f46a3911fa3c0805444483337064	542	1	86.22	4	4
3	0000f6ccb0745a6a4b88665a16c9f078	326	1	43.62	3	4
4	0004aac84e0df4da2b147fca70cf8255	293	1	196.89	3	4

Frequency label breakdown:

1 - These are the most frequent customers. (Frequency > 3)

- 2 - These are the frequent frequent customers. (Frequency = 3)
- 3 - These are the somewhat frequent customers. (Frequency = 2)
- 4 - These are the least frequent customers. (Frequency = 1)

Monetary Labels:

```
ll_m = rfm.Monetary.quantile(0.25)
mid_m = rfm.Monetary.quantile(0.50)
ul_m = rfm.Monetary.quantile(0.75)
print(ll_m, mid_m, ul_m)
```

```
64.0 113.03 203.39
```

```
def monetary_label(money):
    if money <= ll_m:
        return 4
    elif (money > ll_m) and (money <= mid_m):
        return 3
    elif (money > mid_m) and (money <= ul_m):
        return 2
    elif money > ul_m:
        return 1
```

```
rfm['monetary_label'] = rfm.Monetary.apply(monetary_label)
rfm.head()
```

	customer_unique_id	Recency	Frequency	Monetary	recency_label	frequency_label	monetary_label
0	0000366f3b9a7992bf8c76cfd3221e2	116	1	141.90	1	4	2
1	0000b849f77a49e4a4ce2b2a4ca5be3f	119	1	27.19	1	4	4
2	0000f46a3911fa3c0805444483337064	542	1	86.22	4	4	3
3	0000f6ccb0745a6a4b88665a16c9f078	326	1	43.62	3	4	4
4	0004aac84e0df4da2b147fca70cf8255	293	1	196.89	3	4	2

Monetary label breakdown:

- 1 - These are the customers who spend large amount. (Monetary value within the 25% quantile)
- 2 - These are the customers who spend good amount. (Monetary value between 25% and 50% quantile)
- 3 - These are the customers who spend moderately. (Monetary value between 50% and 75% quantile)
- 4 - These are the customers who spend the least. (Monetary value more than 75% quantile)

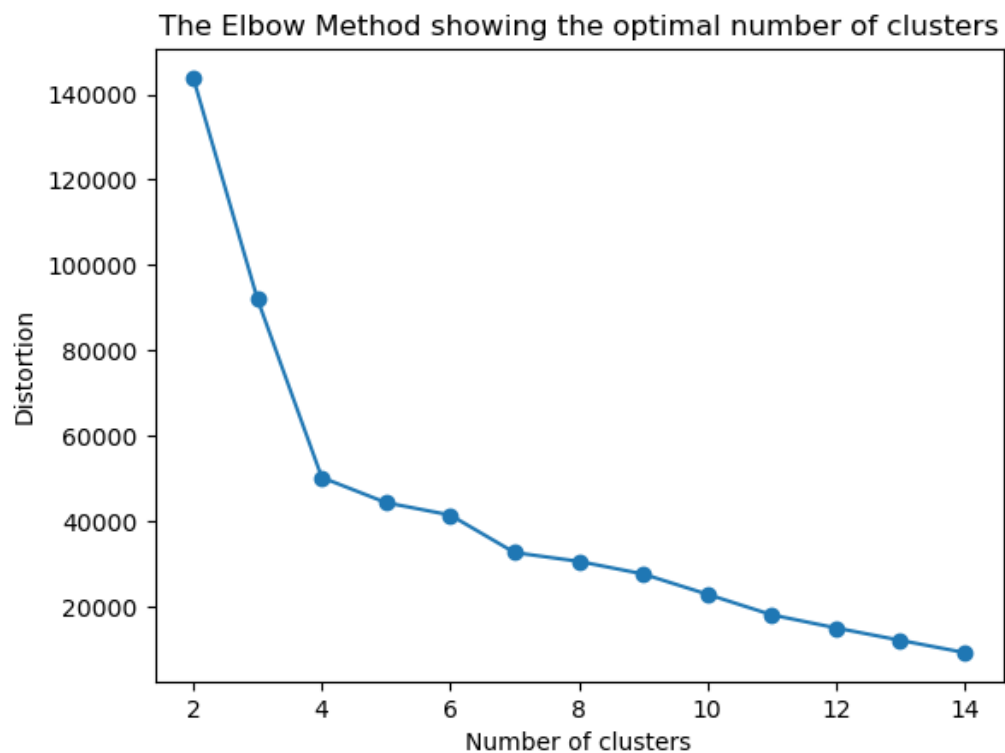
Elbow Method to find optimal number of clusters:

```
: rfm_clstr = rfm.drop(["customer_unique_id", "Recency", "Frequency", "Monetary"], axis='columns')

: from sklearn.cluster import KMeans

distortions=[]
for i in range(2,15):
    kmodel=KMeans(n_clusters=i,n_init=5, random_state=42)
    kmodel.fit(rfm_clstr)
    distortions.append(kmodel.inertia_) # KMeans inertia = Sum of Squares Errors (SSE)

plt.plot(range(2,15), distortions, marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal number of clusters')
plt.show()
```



The number of clusters found to be optimal from the above chart is 10

Applying k-means with n_clusters = 10

```
In [27]: kmodel=KMeans(n_clusters=10, n_init=5, random_state=42)
         kmodel.fit(rfm_clstr)

Out[27]: KMeans(n_clusters=10, n_init=5, random_state=42)

In [28]: from sklearn.metrics import silhouette_score
         from sklearn.metrics import davies_bouldin_score

         print('Silhouette score for K-Means is: ', silhouette_score(rfm_clstr, kmodel.labels_))
         print('Davies-Bouldin Index for K-Means is: ', davies_bouldin_score(rfm_clstr, kmodel.labels_))

Silhouette score for K-Means is:  0.6005961233624046
Davies-Bouldin Index for K-Means is:  0.6334451535987485
```

We use Intrinsic evaluation measures as do not require ground truth labels.

Silhouette Score measures the between-cluster distance against within-cluster distance. A higher score signifies better-defined clusters. The best value is 1 and the worst value is -1.

The Silhouette score for our clustering results is 0.6005

Davies-Bouldin Index measures the size of clusters against the average distance between clusters. A lower score signifies better-defined clusters. The Davies-Bouldin Index (DBI) has the lowest possible value of 0 and does not have an upper limit.

The Davies-Bouldin Index for our clustering results is 0.6334


Overall indicating well-defined clusters.

Plotting clusters in 3 dimensions:

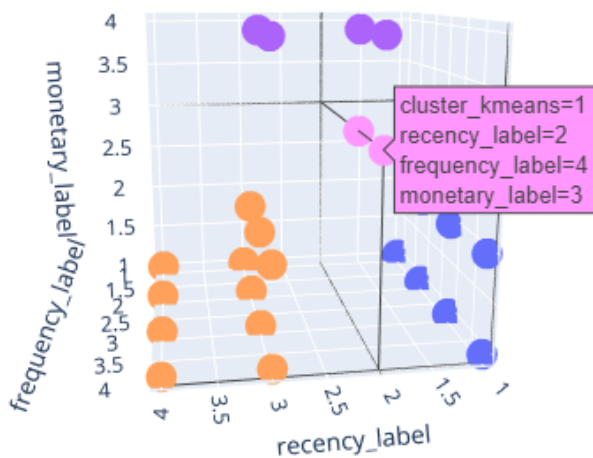
```
: rfm_clstr["cluster_kmeans"] = kmodel.fit_predict(rfm_clstr)

rfm_clstr["cluster_kmeans"] = rfm_clstr["cluster_kmeans"].astype(str)

fig= px.scatter_3d(rfm_clstr, x='recency_label', y='frequency_label', z='monetary_label', color='cluster_kmeans',opacity=1)
fig.update_traces(marker_size = 10)
fig.show()
```



Selecting a certain cluster for better understanding of the visual:



Finding the cluster centers to better understand the clustering results:

```
: kmodel.cluster_centers_  
: array([[3.33146587, 3.93987525, 1.33054176],  
        [1.49761484, 3.98727915, 3.        ],  
        [4.        , 3.98098598, 2.51372279],  
        [1.        , 3.93288702, 1.50880238],  
        [2.49743019, 3.99854377, 4.        ],  
        [4.        , 3.99660268, 4.        ],  
        [2.        , 3.96699897, 2.        ],  
        [1.        , 3.99795327, 4.        ],  
        [3.        , 3.98880348, 3.        ],  
        [2.        , 3.90471276, 1.        ]])
```

CHURN CLASSIFICATION

The churn rate, also known as the rate of attrition or customer churn, is the rate at which customers stop doing business with an entity.

We used the recency column to frame the target variable. If the customer's recency falls above the average value of recency, we consider such customers as churned. The rest of the customers as not churned.

We used the mean of recency as the threshold as the recency is normally or symmetrically distributed.

We will have to impute the target variable to the main dataframe and do the further classification algorithm.

```
: rfm['Churn'] = rfm.Recency.apply(lambda x: 1 if x > rfm.Recency.mean() else 0)
rfm.head()
```

	customer_unique_id	Recency	Frequency	Monetary	recency_label	frequency_label	monetary_label	Churn
0	0000366f3b9a7992bf8c76cfd3221e2	116	1	141.90	1	4	2	0
1	0000b849f77a49e4a4ce2b2a4ca5be3f	119	1	27.19	1	4	4	0
2	0000f46a3911fa3c0805444483337064	542	1	86.22	4	4	3	1
3	0000f6ccb0745a6a4b88665a16c9f078	326	1	43.62	3	4	4	1
4	0004aac84e0df4da2b147fca70cf8255	293	1	196.89	3	4	2	1

Finding certain new calculated attributes

```
|: df['purchased_approved'] = (df.order_approved_at - df.order_purchase_timestamp).dt.seconds
df['approved_carrier'] = (df.order_delivered_carrier_date - df.order_approved_at).dt.days
df['carrier_delivered'] = (df.order_delivered_customer_date - df.order_delivered_carrier_date).dt.days
df['delivered_estimated'] = (df.order_estimated_delivery_date - df.order_delivered_customer_date).dt.days
df['purchased_delivered'] = (df.order_delivered_customer_date - df.order_purchase_timestamp).dt.days
df.head()
```

	Unnamed: 0	customer_id	customer_unique_id	customer_zip_code_prefix	customer_city	customer_state
0	0	06b8999e2fba1a1fbc88172c00ba8bc7	861eff4711a542e4b93843c6dd7febb0	14409	franca	SP 00e7ee1b050b84
1	1	8912fc0c3bbf1e2bf35819e21706718	9eae34bbd3a474ec5d07949ca7de67c0	68030	santarem	PA c1d2b34febe9cc
2	2	8912fc0c3bbf1e2bf35819e21706718	9eae34bbd3a474ec5d07949ca7de67c0	68030	santarem	PA c1d2b34febe9cc
3	3	f0ac8e5a239118859b1734e1087cbb1f	3c799d181c34d51f6d44bbbc563024db	92480	nova santa rita	RS b1a5d5365d330d
4	4	6bc8d08963a135220ed6c6d098831f84	23397e992b09769faf5e66f9e171a241	25931	mage	RJ 2e604b3614664a

5 rows x 44 columns

Grouping by customer id and merging with RFM data

```
In [36]: final = final.merge(rfm[['customer_unique_id', 'Recency', 'Monetary', 'Frequency', 'Churn']], on = 'customer_unique_id')
final.head()
```

Out[36]:

	customer_unique_id	customer_zip_code_prefix	customer_city	customer_state	order_id	purchased_approved	delivered_estimated	purchased_u
0	0000366f3b9a7992bf8c76cfd3221e2	7787	cajamar	SP	1	891.0	4.0	
1	0000b849f77a49e4a4ce2b2a4ca5be3f	6053	osasco	SP	1	26057.0	4.0	
2	0000f46a3911fa3c0805444483337064	88115	sao jose	SC	1	0.0	1.0	
3	0000f6ccb0745a6a4b88665a16c9f078	66812	belem	PA	1	1176.0	11.0	
4	0004aac84e0df4da2b147fca70cf8255	18040	sorocaba	SP	1	1270.0	7.0	

5 rows x 23 columns

The Churn data obtained is balanced but as seen in the previous EDA, certain attributes have a lot of outliers

Outlier Treatment:

```
final_outlierTreated = final.copy()
```

```
for i in final_outlierTreated.select_dtypes(include = np.number).columns:
    q1 = final_outlierTreated[i].quantile(0.25)
    q3 = final_outlierTreated[i].quantile(0.75)
    iqr = q3 - q1
    ul = q3 + 1.5*iqr
    ll = q1 - 1.5*iqr
    final_outlierTreated[i] = np.where(final_outlierTreated[i]>ul,ul,final_outlierTreated[i])
    final_outlierTreated[i] = np.where(final_outlierTreated[i]<ll,ll,final_outlierTreated[i])
```

Grouping states in north south east and west to label encode

```
def state_encoding(state):
    if state in ['RS', 'SC', 'PR']:
        return 'southern'
    elif state in ['SP', 'RJ', 'MG', 'ES']:
        return 'southeastern'
    elif state in ['MT', 'MS', 'GO', 'DF']:
        return 'centralwestern'
    elif state in ['MA', 'PI', 'CE', 'RN', 'PB', 'PE', 'AL', 'SE', 'BA']:
        return 'northeastern'
    else:
        return 'northern'
```

```
final_outlierTreated['customer_state'] = final_outlierTreated['customer_state'].apply(state_encoding)
```

Since the payment_value feature is same as that of the Monetary feature, the former is dropped.

Similarly, customer_city is a multi-class feature, so encoding it would be useless. So we drop the feature.

After deleting Null values, Standard Scaler is applied

```
] : Xi = X.drop(["Churn"], axis = "columns")
```

```
] : from sklearn.preprocessing import StandardScaler  
  
    scaler = StandardScaler().fit(Xi)  
    scaler
```

```
] : StandardScaler()
```

```
] : X_scaled = scaler.transform(Xi)
```

```
] : Y = X['Churn']
```

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X_scaled, Y, random_state = 500, test_size = 0.2)  
  
# check the dimensions of the train & test subset using 'shape'  
# print dimension of train set  
print('X_train', X_train.shape)  
print('y_train', y_train.shape)  
  
# print dimension of test set  
print('X_test', X_test.shape)  
print('y_test', y_test.shape)  
  
X_train (73173, 23)  
y_train (73173,)  
X_test (18294, 23)  
y_test (18294,)
```

Applying Classification Models on Churn Data :

Applying Logistic Regression as it is one of the common Classification models that is robust to outliers.

Logistic Regression

```
In [104]: from sklearn.linear_model import LogisticRegression  
  
log = LogisticRegression(fit_intercept=True, random_state=42, max_iter=2000)  
log.fit(X_train, y_train)
```

```
Out[104]: LogisticRegression(max_iter=2000, random_state=42)
```


```
In [105]: y_pred_log = log.predict(X_test)
```

```
In [106]: from sklearn.metrics import confusion_matrix, classification_report  
from sklearn.metrics import recall_score  
cm_log = confusion_matrix(y_test, y_pred_log)  
print('Confusion Matrix:\n', cm_log)  
print('Classification Report:\n', classification_report(y_test, y_pred_log))  
log_score=log.score(X_test, y_test)  
print(log_score)
```

```
Confusion Matrix:  
[[7603 2498]  
 [4917 3276]]  
Classification Report:  
              precision    recall  f1-score   support  
  
     0       0.61       0.75       0.67       10101  
     1       0.57       0.40       0.47        8193  
  
 accuracy       0.59  
 macro avg       0.59       0.58       0.57       18294  
weighted avg       0.59       0.59       0.58       18294  
  
0.5946758500054663
```

ROC Curve:

```
In [107]: from sklearn.metrics import RocCurveDisplay  
ax = plt.gca()  
plt.plot([0, 1], [0, 1], linestyle='--', label='Random Guessing')  
log_disp = RocCurveDisplay.from_estimator(log, X_test, y_test, ax=ax, alpha=0.8)  
plt.show()
```



Applying various Bagging and Boosting methods to improve accuracy:

Random Forest

```
In [108]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_curve, auc
rfc = RandomForestClassifier(n_estimators=50, random_state=0)
rfc.fit(X_train, y_train)
y_pred_rf = rfc.predict(X_test)
```

```
In [109]: from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import recall_score
cm = confusion_matrix(y_test, y_pred_rf)
print('Confusion Matrix:\n', cm)
print('Classification Report:\n', classification_report(y_test, y_pred_rf))
RFC=rfc.score(X_test, y_test)
print(RFC)
```

Confusion Matrix:

[[8009 2092]

[2788 5405]]

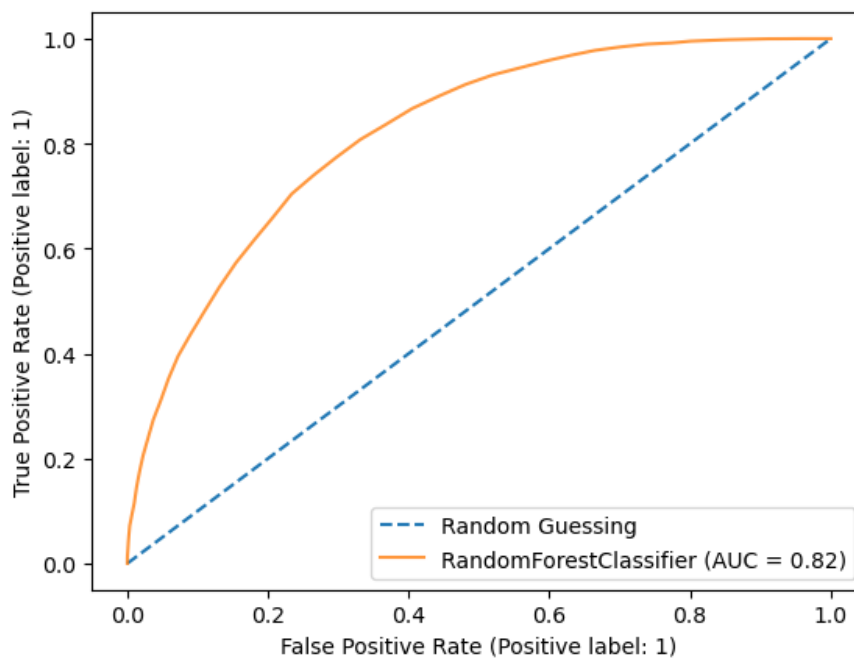
Classification Report:

	precision	recall	f1-score	support
0	0.74	0.79	0.77	10101
1	0.72	0.66	0.69	8193
accuracy			0.73	18294
macro avg	0.73	0.73	0.73	18294
weighted avg	0.73	0.73	0.73	18294

0.7332458729638133

ROC Curve:

```
In [110]: from sklearn.metrics import RocCurveDisplay
ax = plt.gca()
plt.plot([0, 1], [0, 1], linestyle='--', label='Random Guessing')
rfc_disp = RocCurveDisplay.from_estimator(rfc, X_test, y_test, ax=ax, alpha=0.8)
plt.show()
```



AdaBoost

```
In [111]: from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
abc = AdaBoostClassifier(DecisionTreeClassifier(max_depth=2), n_estimators=50, learning_rate=0.1, random_state=0)
abc.fit(X_train, y_train)

y_pred_abc = abc.predict(X_test)
```

```
In [112]: from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import recall_score
cm_ada = confusion_matrix(y_test, y_pred_abc)
print('Confusion Matrix:\n', cm_ada)
print('Classification Report:\n', classification_report(y_test, y_pred_abc))
ada_score=log.score(X_test , y_test)
print(ada_score)
```

Confusion Matrix:
[[7729 2372]
 [3852 4341]]


Classification Report:

	precision	recall	f1-score	support
0	0.67	0.77	0.71	10101
1	0.65	0.53	0.58	8193
accuracy			0.66	18294
macro avg	0.66	0.65	0.65	18294
weighted avg	0.66	0.66	0.65	18294

0.5946758500054663

ROC Curve:

```
In [113]: from sklearn.metrics import RocCurveDisplay
ax = plt.gca()
plt.plot([0, 1], [0, 1], linestyle='--', label='Random Guessing')
ada_disp = RocCurveDisplay.from_estimator(abc, X_test, y_test, ax=ax, alpha=0.8)
plt.show()
```



XGBoost

```
In [114]: from sklearn.metrics import accuracy_score
from xgboost import XGBClassifier
model = XGBClassifier(n_estimators = 50, random_state=0)
model.fit(X_train, y_train)

# Predict on the test set and calculate accuracy
y_pred_xg = model.predict(X_test)
```

```
In [115]: from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import recall_score
cm_xg = confusion_matrix(y_test, y_pred_xg)
print('Confusion Matrix:\n', cm_xg)
print('Classification Report:\n', classification_report(y_test, y_pred_xg))
xg_score = model.score(X_test, y_test)
print(xg_score)
```


Confusion Matrix:
[[7809 2292]
 [2149 6044]]
Classification Report:

	precision	recall	f1-score	support
0	0.78	0.77	0.78	10101
1	0.73	0.74	0.73	8193
accuracy			0.76	18294
macro avg	0.75	0.76	0.75	18294
weighted avg	0.76	0.76	0.76	18294

0.75724281185088

ROC Curve:

```
In [116]: from sklearn.metrics import RocCurveDisplay
ax = plt.gca()
plt.plot([0, 1], [0, 1], linestyle='--', label='Random Guessing')
xg_disp = RocCurveDisplay.from_estimator(model, X_test, y_test, ax=ax, alpha=0.8)
plt.show()
```



Model	Accuracy	AUC of ROC
Logistic Regression	0.5946	0.63
Random Forest (Bagging)	0.7332	0.82
AdaBoost (Boosting)	0.5946	0.73
XGBoost (Boosting)	0.7574	0.85

From the above table we can conclude that the XGBoost Classification Model is the Best choice as it has the highest accuracy score and the highest AUC of ROC .

XGBoost is also Robust to outliers, highly accurate and XGBoost includes regularization techniques that help to prevent overfitting, which is could be a problem in Churn Classification Problems. It uses a combination of L1 and L2 regularization to reduce the complexity of the model, resulting in more robust and accurate predictions.

SALES PREDICTION

Importing merged data and necessary libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
path = './data/'

df = pd.read_csv(path+'merged_data.csv')
```

Finding Sales amounts of Products by using groupby method

```
In [6]: dCat_TopOrders = df.groupby(df['product_id'])['order_id'].nunique().reset_index().sort_values('order_id', ascending = False)
```

Also finding various other attributes like avg_review_score, avg_price, average of estimated_delivery of the product

```
In [36]: cat_reviews = df.groupby(df['product_id'])['review_score'].mean().reset_index().sort_values('review_score', ascending = False)
```

```
In [38]: avg_item_val = df.groupby(df['product_id'])['price'].mean().reset_index().sort_values('price', ascending = False)
```

```
: = df.groupby(df['product_id'])['delivered_estimated'].mean().reset_index().sort_values('delivered_estimated', ascending = False)
```

Merging all DataFrames:

```
In [43]: df_reg = pd.merge(dCat_TopOrders, cat_reviews, on="product_id", how='outer')
df_reg = df_reg.merge(avg_item_val, on="product_id", how='outer')
df_reg = df_reg.merge(est_i_del_time, on="product_id", how='outer')
```

```
In [44]: df_reg
```

Out[44]:

	product_id	order_id	review_score	price	delivered_estimated
0	99a4788cb24856965c36a24e339b6058	458	3.914894	88.175551	9.459566
1	aca2eb7d00ea1a7b8ebd4e68314663af	429	4.020638	71.347655	9.020913
2	422879e10f46682990de24d770e7f83d	351	3.927022	54.827850	9.360947
3	d1c427060a0f73f6b889a5c7c61f2ac4	320	4.096045	137.554802	12.171512
4	389d119b48c3043d311335e499d9c6b	310	4.106173	54.635284	8.920398
...
32166	69ff1e4ad10ccba4d0e1ffc0aa771380	1	3.000000	189.990000	3.000000
32167	69fef0f440d7a4d03f5d883264132dc2	1	5.000000	695.720000	11.000000
32168	69fb24f0cd077f460768e66b89c3565e	1	5.000000	45.000000	8.000000
32169	69faf0d53eb73597d6cfd50175901a56	1	5.000000	29.900000	45.000000
32170	ffe9eeff12fcb74a2f2b007dde0c58	1	4.000000	249.990000	-3.000000

32171 rows x 5 columns

Merging with product_id to get product_category and renaming incorrectly labelled data

```
In [46]: dict = {'order_id': 'sales_amt',
               'review_score': 'avg_review_score',
               'price': 'avg_price',
               'delivered_estimated': 'estimated_delivery',
               }

# call rename () method
df_reg.rename(columns=dict,
              inplace=True)
```

```
In [47]: a = df.loc[:, ['product_id', 'product_category']]
```

```
In [48]: sales = pd.merge(df_reg, a, on="product_id", how="right")
```

```
In [49]: sales
```

Out[49]:

	product_id	sales_amt	avg_review_score	avg_price	estimated_delivery	product_category
0	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	Furniture
1	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	Furniture
2	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	Furniture
3	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	Furniture
4	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	Furniture
...
115604	ea9b0b855335919945731f9368f83dc9	1	5.000000	193.000000	17.000000	Home & Garden
115605	0c800efe70e04fcc3b266946e3e4826	1	4.000000	389.000000	11.000000	Home & Garden
115606	775596b5ab8f1cb5890c7263c1c92bc4	1	5.000000	139.000000	13.000000	Home & Garden
115607	cc9e875c2df286dbed83efe01191162c	1	5.000000	129.000000	18.000000	Home & Garden
115608	cc9e875c2df286dbed83efe01191162c	1	5.000000	129.000000	18.000000	Home & Garden

Label Encoding Product_id and deleting null values

```
In [56]: from sklearn import preprocessing

label_encoder = preprocessing.LabelEncoder()

sales['product_category'] = label_encoder.fit_transform(sales['product_category'])
```

```
In [58]: sales
```

Out[58]:

	product_id	sales_amt	avg_review_score	avg_price	estimated_delivery	product_category
0	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	6
1	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	6
2	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	6
3	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	6
4	a9516a079e37a9c9c36b9b78b10169e8	40	3.393443	119.285082	10.245902	6
...
115604	ea9b0b855335919945731f9368f83dc9	1	5.000000	193.000000	17.000000	7
115605	0c800efe70e04fcc3b266946e3e4826	1	4.000000	389.000000	11.000000	7
115606	775596b5ab8f1cb5890c7263c1c92bc4	1	5.000000	139.000000	13.000000	7
115607	cc9e875c2df286dbed83efe01191162c	1	5.000000	129.000000	18.000000	7
115608	cc9e875c2df286dbed83efe01191162c	1	5.000000	129.000000	18.000000	7

115609 rows x 6 columns

```
In [59]: sales.isnull().sum()
```

Out[59]:

product_id	0
sales_amt	0
avg_review_score	0
avg_price	0
estimated_delivery	935
product_category	0
dtype:	int64

```
In [61]: sales.dropna(axis="rows", inplace=True)
```

Train Test Split

```
In [64]: from sklearn.model_selection import train_test_split

x = sales.drop(['product_id', 'sales_amt'], axis='columns')
y = sales[['sales_amt']]
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 1/3, random_state = 101)
```

Applying Regression Models:

Applying Linear Regression as it is the simplest and most common Regression algorithms

Linear Regression

```
In [65]: from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(x_train,y_train)
```

```
Out[65]: LinearRegression()
```

```
In [66]: y_pred_linear = regressor.predict(x_test)
```

```
In [67]: from sklearn import metrics  
  
print("R squared value",metrics.r2_score(y_test,y_pred_linear))  
print("MAE",metrics.mean_absolute_error(y_test,y_pred_linear))  
print("MSE",np.sqrt(metrics.mean_squared_error(y_test,y_pred_linear)))
```

```
R squared value 0.005427260737199013  
MAE 39.14167447925617  
MSE 68.9350864594877
```

Applying Tree based Regressors as they have high explainability thus you can explain the decisions, identify possible events that might occur, and see potential outcomes. The analysis helps you determine what the best decision would be.

High explainability can also be a huge advantage when working with models that will be making business decisions so as to better understand how and why certain decisions are made

DecisionTree Regression

```
In [94]: from sklearn.tree import DecisionTreeRegressor

In [95]: dt = DecisionTreeRegressor(max_depth=10)

In [96]: dt.fit(x_train,y_train)

Out[96]: DecisionTreeRegressor(max_depth=10)

In [97]: y_pred_dt = dt.predict(x_test)

In [98]: print("R squared value",metrics.r2_score(y_test,y_pred_dt))
print("MAE",metrics.mean_absolute_error(y_test,y_pred_dt))
print("MSE",np.sqrt(metrics.mean_squared_error(y_test,y_pred_dt)))

R squared value 0.8925897392034384
MAE 9.418919503661787
MSE 22.6540020005456
```

RandomForestRegressor

```
In [99]: from sklearn.ensemble import RandomForestRegressor

rf= RandomForestRegressor(max_depth=10, min_samples_split = 5)

In [100]: y_train = np.ravel(y_train)

In [101]: rf.fit(x_train,y_train)

Out[101]: RandomForestRegressor(max_depth=10, min_samples_split=5)

In [102]: y_pred_rf= rf.predict(x_test)

In [103]: print("R squared value",metrics.r2_score(y_test,y_pred_rf))
print("MAE",metrics.mean_absolute_error(y_test,y_pred_rf))
print("MSE",np.sqrt(metrics.mean_squared_error(y_test,y_pred_rf)))

R squared value 0.913615379253922
MAE 9.01909850531045
MSE 20.316096379170787
```

XGBoostRegressor

```
In [104]: from xgboost import XGBRegressor

In [109]: xgb=XGBRegressor(n_estimators=100,random_state=42, learning_rate=0.05)

In [110]: xgb.fit(x_train,y_train)

Out[110]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                        colsample_bylevel=None, colsample_bynode=None,
                        colsample_bytree=None, early_stopping_rounds=None,
                        enable_categorical=False, eval_metric=None, feature_types=None,
                        gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
                        interaction_constraints=None, learning_rate=0.05, max_bin=None,
                        max_cat_threshold=None, max_cat_to_onehot=None,
                        max_delta_step=None, max_depth=None, max_leaves=None,
                        min_child_weight=None, missing=nan, monotone_constraints=None,
                        n_estimators=100, n_jobs=None, num_parallel_tree=None,
                        predictor=None, random_state=42, ...)

In [111]: y_pred_xgb = xgb.predict(x_test)

In [112]: print("R squared value",metrics.r2_score(y_test,y_pred_xgb))
          print("MAE",metrics.mean_absolute_error(y_test,y_pred_xgb))
          print("MSE",np.sqrt(metrics.mean_squared_error(y_test,y_pred_xgb)))

R squared value 0.9459609561509599
MAE 10.262075330105832
MSE 16.068519977477678
```

Model	R-squared value	MAE	MSE
Linear Regression	0.00542	39.14	68.93
DecisionTree Regression	0.8925	9.41	20.31
RandomForestRegressor	0.9136	9.02	20.32
XGBoostRegressor	0.9459	10.26	16.06

From the Above table we can observe that the XGBoost Regressor has the highest value of R-square score and has the lowest value of Mean Absolute Error and Mean Squared Error

This is possibly because, XGBoost is a Boosting method that uses several weak learners and improves on the weak learners in every iteration to improve its results.